

### data descriptions:

- id: listing id
- review\_scores\_location: 0-5 stars converted into a 0-10 scale
- name: listing name
- host\_id: host id
- host\_name: host name
- neighbourhood\_group: NYC borough
- neighbourhood: NYC neighborhood
- latitude: listing latitude
- longitude: listing longitude
- room\_type: type of listing (Entire home/apt, Private room, Shared room)
- price: listing price
- minimum\_nights: required minimum nights stay
- number\_of\_reviews: total number of reviews
- last\_review: date of last review
- reviews\_per\_month: average number of reviews per month
- calculated\_host\_listings\_count: total number of listings for this host
- availability\_365: number of days listing is available out of 365

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import numpy as np
import matplotlib.image as mpimg
import geopandas
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import GridSearchCV
```

```
In [2]: ori_data=pd.read_csv("AB_NYC_2019.csv")
```


```
In [3]: print(f'total rows:{ori_data.shape[0]}')
print(f'total columns:{ori_data.shape[1]}')
print(f'column name:{ori_data.columns.values}')
```

```
total rows:48895
total columns:16
column name:['id' 'name' 'host_id' 'host_name' 'neighbourhood_group' 'neighbourhood'
'latitude' 'longitude' 'room_type' 'price' 'minimum_nights'
'number_of_reviews' 'last_review' 'reviews_per_month'
'calculated_host_listings_count' 'availability_365']
```


```
In [4]: ## which columns has duplicated values
ori_data.apply(lambda x:x.unique().shape[0],axis=0)
```

```
Out[4]: id                48895
name                47906
host_id            37457
host_name          11453
neighbourhood_group    5
neighbourhood         221
latitude            19048
longitude           14718
room_type            3
price                674
minimum_nights        109
number_of_reviews     394
last_review          1765
reviews_per_month     938
calculated_host_listings_count  47
availability_365      366
dtype: int64
```

```
In [5]: ## The lower y ,the highly repeat for x (may be categorical variable)
## With high y,x may be continuous variable or index
(ori_data.apply(lambda x:x.unique().shape[0],axis=0)/ori_data.shape[0]).plot()
plt.show()
```

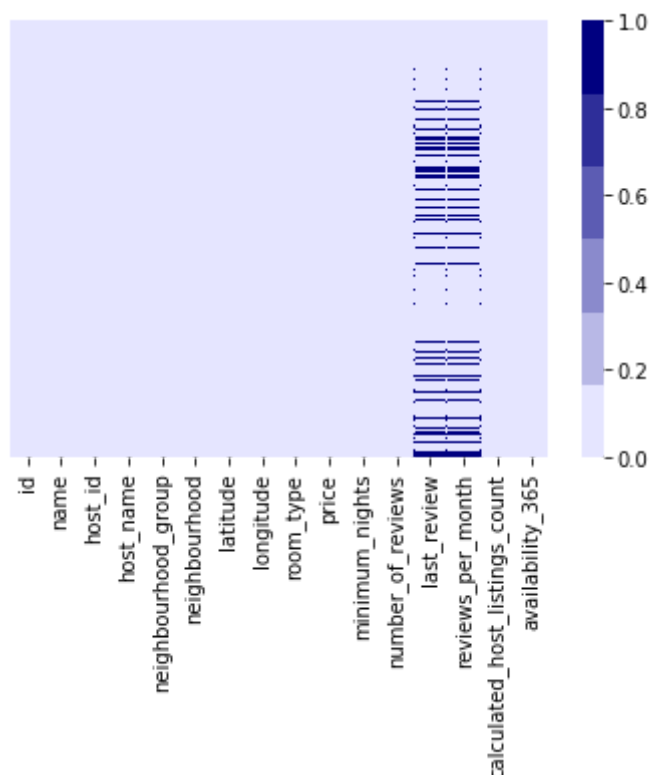


```
In [6]: ##which column contains NaN
ori_data.isna().sum().plot(kind='bar',title='number of missing values')
plt.show()
ori_data.isna().sum()
```



```
Out[6]: id          0
        name        16
        host_id     0
        host_name   21
        neighbourhood_group  0
        neighbourhood  0
        latitude    0
        longitude   0
        room_type   0
        price       0
        minimum_nights  0
        number_of_reviews  0
        last_review  10052
        reviews_per_month  10052
        calculated_host_listings_count  0
        availability_365  0
        dtype: int64
```

```
In [7]: cmap=sns.light_palette('navy',reverse=False)
sns.heatmap(ori_data.isnull().astype(np.int8),yticklabels=False,cmap=cmap)
plt.show()
```



The goal is to predict the price by the geographical and accommodation\_related features. So, we should not take into account host-name and place-name. However, we also can find out that there are many missing data in 'last\_review' and 'reviews\_per\_month'. If the place is a new post, it would not have any review before. Thus, we can add a feature determining whether it is a new post or not.

```
In [8]: ori_data.describe()[['price', 'minimum_nights', 'number_of_reviews', 'reviews_per_month', 'calculated_host_listings_count', 'availability_365']]
```

```
Out[8]:
```

|       | price        | minimum_nights | number_of_reviews | reviews_per_month | calculated_host_listings_count | availability_365 |
|-------|--------------|----------------|-------------------|-------------------|--------------------------------|------------------|
| count | 48895.000000 | 48895.000000   | 48895.000000      | 38843.000000      | 48895.000000                   | 48895.000000     |
| mean  | 152.720687   | 7.029962       | 23.274466         | 1.373221          | 1.000000                       | 161.996000       |
| std   | 240.154170   | 20.510550      | 44.550582         | 1.680442          | 0.000000                       | 149.999999       |
| min   | 0.000000     | 1.000000       | 0.000000          | 0.010000          | 0.000000                       | 0.000000         |
| 25%   | 69.000000    | 1.000000       | 1.000000          | 0.190000          | 0.000000                       | 30.000000        |
| 50%   | 106.000000   | 3.000000       | 5.000000          | 0.720000          | 0.000000                       | 60.000000        |
| 75%   | 175.000000   | 5.000000       | 24.000000         | 2.020000          | 0.000000                       | 90.000000        |
| max   | 10000.000000 | 1250.000000    | 629.000000        | 58.500000         | 0.000000                       | 365.000000       |

```
In [9]: for col in ['price', 'minimum_nights', 'number_of_reviews', 'reviews_per_month']:  
        print(f'{col}:{ori_data[col].quantile(0.95) }')
```

```
price:355.0  
minimum_nights:30.0  
number_of_reviews:114.0  
reviews_per_month:4.64  
calculated_host_listings_count:15.0  
availability_365:359.0
```

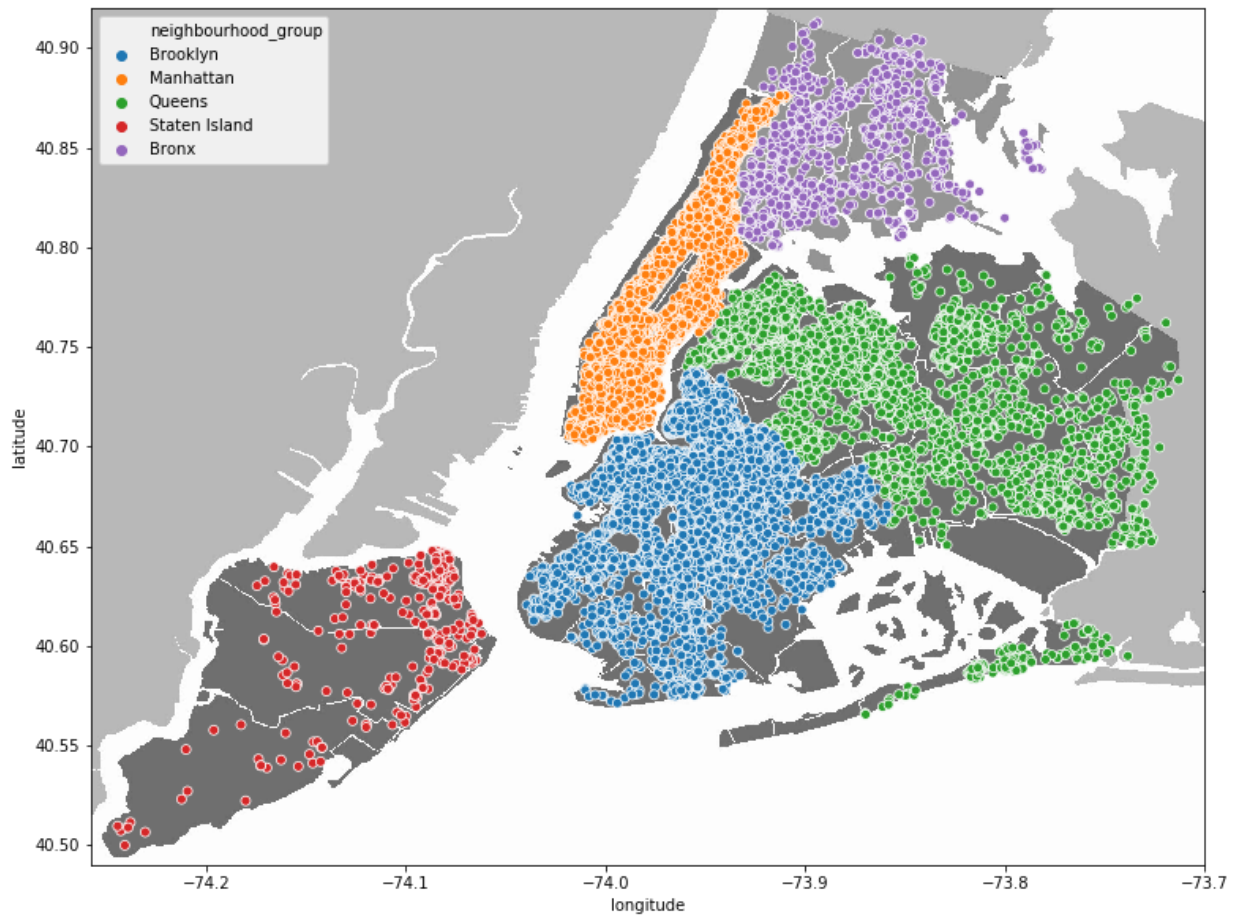
The outliers in "price" and "minimum\_nights" are unreasonable. Thus, we can filter out top 5% outlier in this two columns.

```
In [10]: price_filter=(ori_data['price']<ori_data['price'].quantile(0.95))&(ori_data  
minimum_nights_filter=ori_data['minimum_nights']<ori_data['minimum_nights']  
filter_data=ori_data[price_filter&minimum_nights_filter].copy()
```

```
In [11]: print(f'original_data:{ori_data.shape[0]}')  
        print(f'filtered_data:{filter_data.shape[0]}')
```


```
original_data:48895  
filtered_data:42201
```

```
In [12]: plt.figure(figsize=(20,10))
nyc_img = plt.imread("./New_York_City.png",0)
plt.imshow(nyc_img,zorder=0,extent=[-74.258, -73.7, 40.49,40.92])
ax=plt.gca()
sns.scatterplot(x=filter_data.longitude,y=filter_data.latitude,hue=ori_data)
plt.show()
```



We compare price based on location on the map. Manhattan gathers most high-priced places.

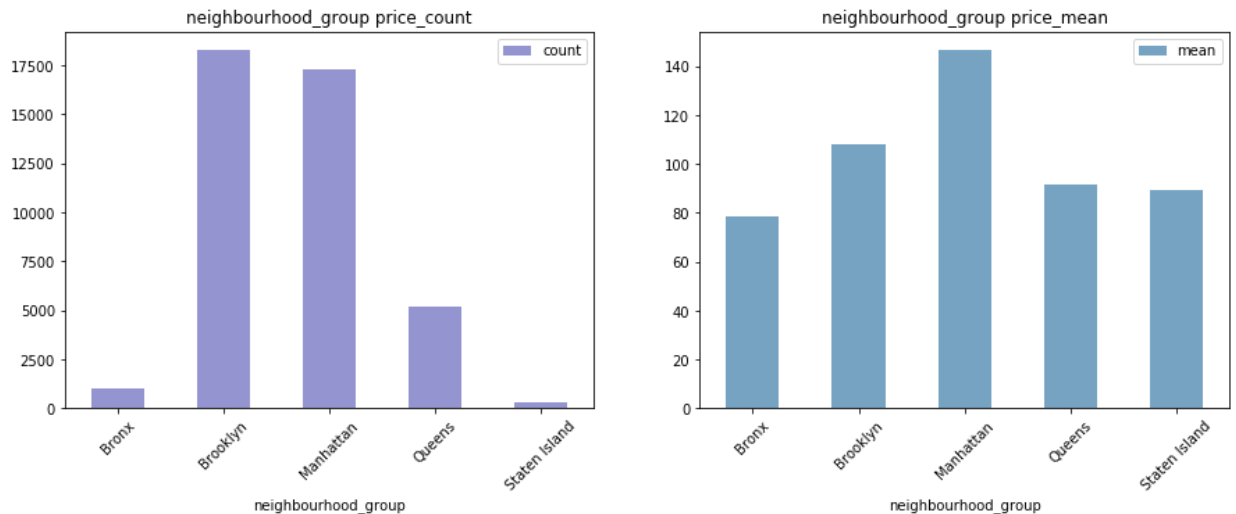
```
In [13]: plt.figure(figsize=(20,10))
nyc_img = plt.imread("./New_York_City.png",0)
ax=plt.gca()
##filter outlier
plt.imshow(nyc_img,zorder=0,extent=[-74.258, -73.7, 40.49,40.92])
sc=plt.scatter(filter_data.longitude, filter_data.latitude, c=filter_data.p)
plt.colorbar(sc)
plt.show()
```



```
In [14]: ##only 5 unique values in neighbourhood_group
filter_data.neighbourhood_group.unique()
```

```
Out[14]: array(['Brooklyn', 'Manhattan', 'Queens', 'Staten Island', 'Bronx'],
              dtype=object)
```


```
In [15]: neighbourhood_group=filter_data.groupby("neighbourhood_group")["price"].des
fig, axes = plt.subplots(nrows=1, ncols=2,figsize=(15,5))
neighbourhood_group.plot(x="neighbourhood_group", y=["count"], kind="bar",rc
neighbourhood_group.plot(x="neighbourhood_group", y=["mean"], kind="bar",rc
plt.show()
```



- More than 17000 posts in Manhattan and Brooklyn.
- The average price in Manhattan is most expensive and less in Bronx.




```
In [16]: plt.figure(figsize=(10,5))
filter_data.groupby('room_type')['price'].mean().plot( kind="bar",rot=0,col
plt.xlabel('room_type',fontsize=12)
plt.ylabel('price',fontsize=12)
plt.title("bar chart(room_type & price)",fontsize=16)
plt.show()
```



- The average price for Entire home/apt is higher than Private room and Shared room.

```
In [17]: plt.figure(figsize=(10,10))
sns.catplot(x='neighbourhood_group',y='price',data=filter_data,hue='room_ty
plt.show()
```

<Figure size 720x720 with 0 Axes>



- Most area in New York, the average price of Shared room is far lower than other room\_type. However, the price in Staten Island is close to Private room.


```

In [18]: #number_of_reviews*reviews_per_month=post_month(how long did the place be p
filter_data['posted_month']=round(filter_data['number_of_reviews']/filter_c

In [19]: filter_data['is_New']=filter_data['reviews_per_month'].apply(lambda x:1 if


In [20]: plt.figure(figsize=(10,5))
sns.scatterplot(x=filter_data.number_of_reviews,y=filter_data.price,hue=fil
plt.show()

```



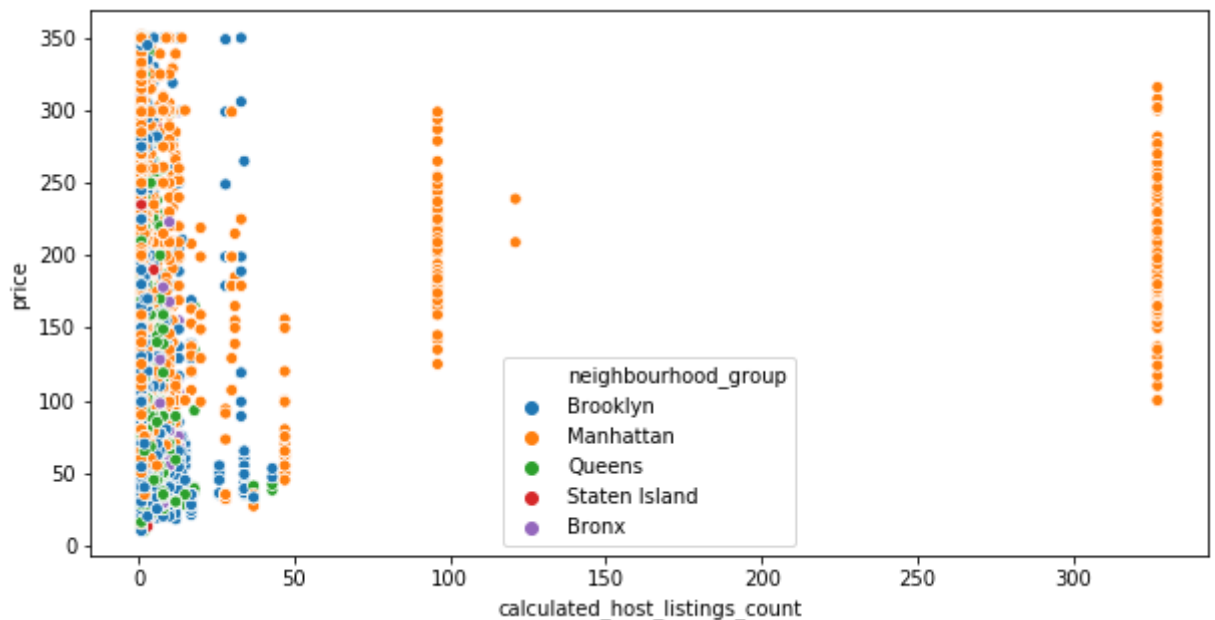
- Reviews more than 400 are gather in Private room.
- Though the number of places in Queens is not highest but several places in Queens had more than 400 reviews.

```
In [21]: plt.figure(figsize=(10,5))
sns.scatterplot(x=filter_data.number_of_reviews,y=filter_data.price,hue=filter_data.neighbourhood_group)
plt.show()
```




- Most of hosts had less than 50 host listings and few hosts in Manhattan had more than 50 host listings.

```
In [22]: plt.figure(figsize=(10,5))
sns.scatterplot(x=filter_data.calculated_host_listings_count,y=filter_data.price)
plt.show()
```



```
In [112]: filter_data.groupby('is_New')['price'].count().plot( kind="bar",rot=0,color
plt.ylabel("count")
```

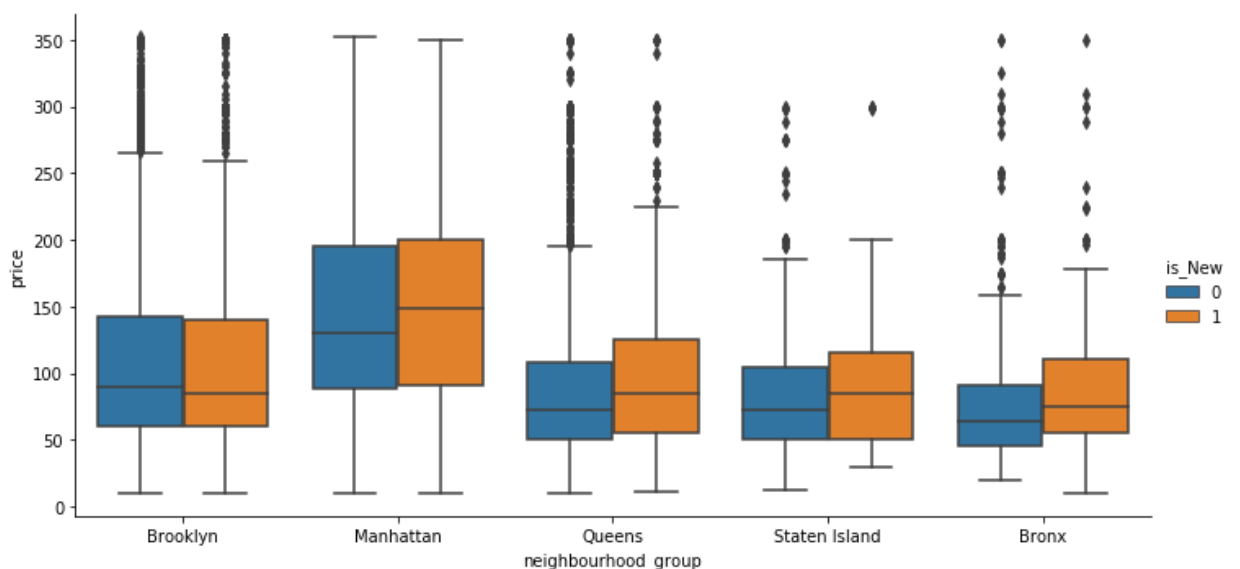
```
Out[112]: Text(0, 0.5, 'count')
```



In Brooklyn ,the average price of new-posted places is less than that of old places but in contrast with other areas.

```
In [24]: plt.figure(figsize=(10,10))
sns.catplot(x='neighbourhood_group',y='price',data=filter_data,hue='is_New'
plt.show()
```


<Figure size 720x720 with 0 Axes>



- For room-type, we can see the average price for all type of new-posted places had higher price than before.

```
In [25]: plt.figure(figsize=(10,10))
sns.catplot(x='room_type',y='price',data=filter_data,hue='is_New', kind="box")
plt.show()
```

<Figure size 720x720 with 0 Axes>

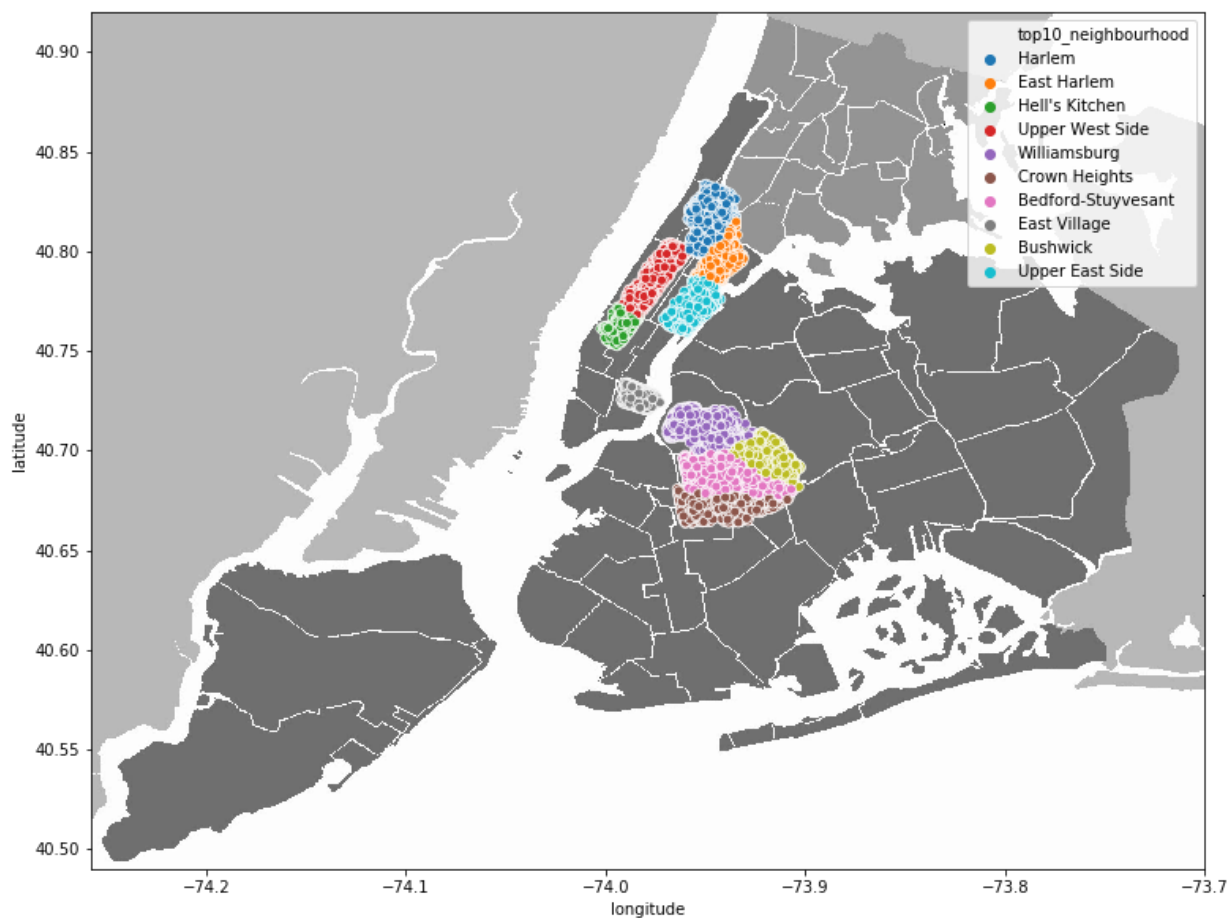


### the list of top 10 neighbourhood in NYC

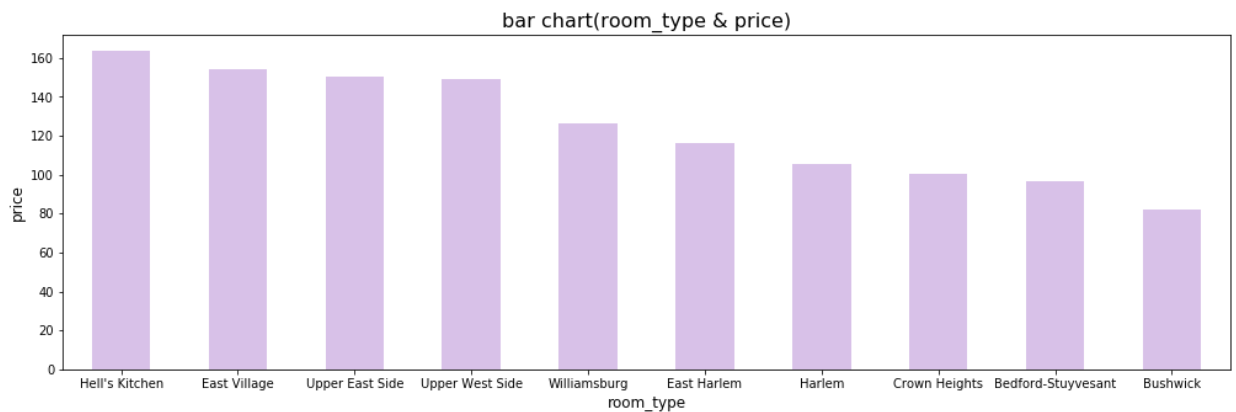
```
In [26]: #neighbourhood value counts
top10_list=filter_data.neighbourhood.value_counts()[:10].index
def top10(x):
    if(x in top10_list):
        return x
filter_data['top10_neighbourhood']=filter_data['neighbourhood'].apply(top10)
```

- Top-10-favored neighbourhood are gathered in Manhattan and Brooklyn.

```
In [27]: plt.figure(figsize=(20,10))
nyc_img = plt.imread("./New_York_City.png",0)
plt.imshow(nyc_img,zorder=0,extent=[-74.258, -73.7, 40.49,40.92])
ax=plt.gca()
sns.scatterplot(x=filter_data.longitude,y=filter_data.latitude,hue=filter_c
plt.show()
```



```
In [28]: plt.figure(figsize=(17,5))
filter_data.groupby('top10_neighbourhood')['price'].mean().sort_values(ascending=True)
plt.xlabel('room_type',fontsize=12)
plt.ylabel('price',fontsize=12)
plt.title("bar chart(room_type & price)",fontsize=16)
plt.show()
```



Look into "top10\_neighbourhood", the chart above shows the average price of places near by most popular neighbourhood. The top-4 leaders are all in Manhattan.

```
In [29]: filter_data=filter_data.drop(columns=['neighbourhood','id','name','host_id'])
```

```
In [30]: ##fillna with zero
filter_data=filter_data.fillna('0')
filter_data=filter_data[['neighbourhood_group','latitude','longitude','room_type',
'minimum_nights','number_of_reviews','reviews_per_month',
'calculated_host_listings_count','availability_365','posted_month',
'is_New','top10_neighbourhood','price']]
```

```
In [31]: backup=filter_data.copy()
#filter_data=backup.copy()
```

```
In [32]: ori_price=backup[['price']]
```

```
In [33]: filter_data.columns
```

```
Out[33]: Index(['neighbourhood_group', 'latitude', 'longitude', 'room_type',
'minimum_nights', 'number_of_reviews', 'reviews_per_month',
'calculated_host_listings_count', 'availability_365', 'posted_month',
'is_New', 'top10_neighbourhood', 'price'],
dtype='object')
```

```
In [34]: ##label encoding for categorical features
##StandardScaler for continuous features
continous_columns=['latitude', 'longitude', 'minimum_nights', 'number_of_re
            'reviews_per_month', 'calculated_host_listings_count', 'availability_
categorical_columns=['neighbourhood_group', 'room_type', 'is_New', 'top10_nei
labelencoder = LabelEncoder()
for col in categorical_columns:
    filter_data[col] = labelencoder.fit_transform(filter_data[col])
```

```
In [35]: scaler = StandardScaler()
scaler.fit(filter_data[continous_columns])
filter_data[continous_columns] = scaler.transform(filter_data[continous_col
```

```
In [36]: ##train_test_split
y=filter_data[['price']]
X=filter_data.drop(columns=['price'])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
```

```
In [37]: regr = RandomForestRegressor()
regr.fit(X_train, y_train)
y_predict=regr.predict(X_test)
mse = np.mean((regr.predict(X_test)-y_test['price']) ** 2)
r_squared=regr.score(X_test, y_test)
adj_r_squared = r_squared - (1 - r_squared) * (X_test.shape[1] / (X_test.sh
print(f"MSE:{mse}")
print(f"r_squared:{r_squared}")
print(f"adj_r_squared:{adj_r_squared}")
```

/usr/local/lib/python3.7/site-packages/sklearn/ensemble/forest.py:245: FutureWarning: The default value of n\_estimators will change from 10 in version 0.20 to 100 in 0.22.

"10 in version 0.20 to 100 in 0.22.", FutureWarning)

/usr/local/lib/python3.7/site-packages/ipykernel\_launcher.py:2: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```
MSE:0.48889020985651976
r_squared:0.5201393449254608
adj_r_squared:0.5194561071631335
```

## Tuning Parameters



```
In [40]: # Number of trees in random forest
n_estimators = [10,20,40,50,100,200,400,800,1000]
# Number of features to consider at every split
max_features = ['auto', 'sqrt']
# Maximum number of levels in tree
max_depth = [2,5,8,10,12,20,50,80]
max_depth.append(None)
# Minimum number of samples required to split a node
min_samples_split = [2, 5, 10]
# Minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 4]
# Method of selecting samples for training each tree
bootstrap = [True, False]
# Create the random grid
tuned_parameters = {'n_estimators': n_estimators,
                    'max_features': max_features,
                    'max_depth': max_depth,
                    'min_samples_split': min_samples_split,
                    'min_samples_leaf': min_samples_leaf,
                    'bootstrap': bootstrap}
print(tuned_parameters)
```

```
{'n_estimators': [10, 20, 40, 50, 100, 200, 400, 800, 1000], 'max_features': ['auto', 'sqrt'], 'max_depth': [2, 5, 8, 10, 12, 20, 50, 80, None], 'min_samples_split': [2, 5, 10], 'min_samples_leaf': [1, 2, 4], 'bootstrap': [True, False]}
```

```
In [41]: rf = RandomForestRegressor()
         clf = GridSearchCV(rf, tuned_parameters, n_jobs=-1, verbose=1)
         clf.fit(X_train, y_train)
```

Fitting 3 folds for each of 2916 candidates, totalling 8748 fits

/home/jim/.local/lib/python3.6/site-packages/sklearn/model\_selection/\_split.py:1978: FutureWarning: The default value of cv will change from 3 to 5 in version 0.22. Specify it explicitly to silence this warning.

warnings.warn(CV\_WARNING, FutureWarning)

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 16 concurrent workers.

|  |                            |
|--|----------------------------|
| [Parallel(n_jobs=-1)]: Done 18 tasks         | elapsed: 3.1s              |
| [Parallel(n_jobs=-1)]: Done 168 tasks        | elapsed: 49.1s             |
| [Parallel(n_jobs=-1)]: Done 418 tasks        | elapsed: 1.6min            |
| [Parallel(n_jobs=-1)]: Done 768 tasks        | elapsed: 4.8min            |
| [Parallel(n_jobs=-1)]: Done 1218 tasks       | elapsed: 9.8min            |
| [Parallel(n_jobs=-1)]: Done 1768 tasks       | elapsed: 17.4min           |
| [Parallel(n_jobs=-1)]: Done 2418 tasks       | elapsed: 26.8min           |
| [Parallel(n_jobs=-1)]: Done 3168 tasks       | elapsed: 46.1min           |
| [Parallel(n_jobs=-1)]: Done 4018 tasks       | elapsed: 64.9min           |
| [Parallel(n_jobs=-1)]: Done 4968 tasks       | elapsed: 75.6min           |
| [Parallel(n_jobs=-1)]: Done 6018 tasks       | elapsed: 93.9min           |
| [Parallel(n_jobs=-1)]: Done 7168 tasks       | elapsed: 125.5min          |
| [Parallel(n_jobs=-1)]: Done 8418 tasks       | elapsed: 171.3min          |
| [Parallel(n_jobs=-1)]: Done 8748 out of 8748 | elapsed: 180.8min finished |

/home/jim/.local/lib/python3.6/site-packages/sklearn/model\_selection/\_search.py:715: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```
self.best_estimator_.fit(X, y, **fit_params)
```

```
Out[41]: GridSearchCV(cv='warn', error_score='raise-deprecating',
                    estimator=RandomForestRegressor(bootstrap=True, criterion='mse',
                                                    max_depth=None,
                                                    max_features='auto',
                                                    max_leaf_nodes=None,
                                                    min_impurity_decrease=0.0,
                                                    min_impurity_split=None,
                                                    min_samples_leaf=1,
                                                    min_samples_split=2,
                                                    min_weight_fraction_leaf=0.0,
                                                    n_estimators='warn', n_jobs=None,
                                                    oob_score=False, random_state=None,
                                                    verbose=0, warm_start=False),
                    iid='warn', n_jobs=-1,
                    param_grid={'bootstrap': [True, False],
                                'max_depth': [2, 5, 8, 10, 12, 20, 50, 80, None],
                                'max_features': ['auto', 'sqrt'],
                                'min_samples_leaf': [1, 2, 4],
                                'min_samples_split': [2, 5, 10],
```

```

        'n_estimators': [10, 20, 40, 50, 100, 200, 400,
800,
                                1000]},
        pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
        scoring=None, verbose=1)

```

```
In [42]: print(clf.best_params_)
```

```

{'bootstrap': True, 'max_depth': 20, 'max_features': 'sqrt', 'min_samples_
_leaf': 1, 'min_samples_split': 10, 'n_estimators': 1000}

```

```
In [39]: best_params_={'bootstrap': True, 'max_depth': 20, 'max_features': 'sqrt', 'n_estimators': 1000}
```

```
In [40]: regr = RandomForestRegressor(**best_params_)
regr.fit(X_train, y_train)
y_predict=regr.predict(X_test)
mse = np.mean((y_predict-y_test['price']) ** 2)
r_squared=regr.score(X_test, y_test)
adj_r_squared = r_squared - (1 - r_squared) * (X_test.shape[1] / (X_test.shape[1] + 1))
print(f"MSE:{mse}")
print(f"r_squared:{r_squared}")
print(f"adj_r_squared:{adj_r_squared}")

```

/usr/local/lib/python3.7/site-packages/ipykernel\_launcher.py:2: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```


MSE:0.43945942074606587
r_squared:0.5686571723746036
adj_r_squared:0.5680430155246385

```

```
In [82]: test_df=X_test.copy()
test_df['price']=y_predict
test_df[continuous_columns]=scaler.inverse_transform(test_df[continuous_columns])
test_df=test_df.reset_index()
real_df=ori_price.reset_index().rename(columns={'price':'real_price'})
test_comparsion=pd.merge(test_df,real_df,how='left')[['price','real_price']]
test_comparsion=test_comparsion.sort_values('price').reset_index(drop=True)

```

```
In [106]: x=list(test_comparsion.index[1:21])
y=list(test_comparsion['price'].values[1:21])
y2=list(test_comparsion['real_price'].values[1:21])
fig, ax = plt.subplots(figsize=(20, 5))
ax.scatter(x, y,label='predict')
ax.scatter(x, y2,label='real')
plt.title('First 20 Places Price Prediction in test dataset',fontsize=30)
plt.yticks(np.linspace(0,350,8),fontsize=15)
plt.xticks(np.linspace(1,20,20),fontsize=15)
plt.xlabel('Place',fontsize=20)
plt.ylabel('Price',fontsize=20)
plt.legend(loc='best', fontsize = 20)
plt.show()
```



```
In [ ]:
```